

Job Scheduler 3.0 Component Specification

1. Design

The Job Scheduling Component enables the timed execution of specified tasks. This functionality is similar to the unix cron utility and variants that exist on most operating system. Users can schedule both one-time and repeating tasks.

In the second version of this component, the concepts of Job dependency and grouping were introduced to create a more robust capability. Jobs could be created that are dependent on other jobs, and to more easily perform jobs by grouping them. The third version, as proposed by this design, is to separate the scheduling part of the design from the processing part. This new scheduler will be solely responsible for managing the jobs, including their persistence. The processor will be moved to a separate component.

Versions 1.0 and 2.0

A job is a specific O/S Command/java class to be launched on a particular schedule. The Job information is stored in a configuration file. Jobs can be created either through this component or through manual edit of the configuration file.

A job can be of two specific types. The first – External – are for operating system executables and scripts. This component uses the Executable Wrapper to execute these. The second – Internal – are Java classes implementing both the Runnable and Schedulable interfaces, the later being in this component. The component uses reflection to load and execute these classes.

Support for multiple schedules is provided by re-adding a job with a different name and a new schedule.

Both internal and external jobs are executed asynchronously. This permits simultaneous and overlapping jobs.

The Logging Wrapper component is used to log the results (run start, run completion, status) of each launched job. This provides an essential audit trail of all execution attempts and results.

A job can be dependent on another job regarding execution time. For example Job B does not have a configured date time. It is configured to run on the successful completion of Job A. Allow the option to execute Job B even if Job A fails. Also allow configuration of an option time delay before Job B executes after Job A. However, a Job must have either a scheduled date time or a task dependency but not both.



The component allows the set up of an email alert if a job returns a failure or cannot be executed. Configuration is at the job or job group level.

A job group is a grouping of jobs to simplify job scheduler configuration. A group can contain one or more jobs. A job can belong to one or more groups. The job groups are used to ease “Alert Notification Configuration”.

Versions 3.0

Version 3.0 will incorporate the following functionality:

- 1) The processing will be removed from the Scheduler class and moved to a separate component.
- 2) The scheduling of a job will be enhanced beyond the use of Calendars and simple intervals.
- 3) Giving the scheduler CRUD functionality.
- 4) Retaining as much of the current design

1) Refactoring

The first point will be fairly straightforward. The Current design already has a separate processor entity. What the new design will do is refactor all Scheduler tasks that deal with running the jobs strictly to the processor, and strictly deal with the persistence of the Jobs. As such, the current relationship between the Scheduler and the Processor will be reversed: Now, the processor will query the scheduler for Jobs, instead of the Scheduler feeding the Jobs to the Processor as they are added.

2) Scheduling

The second point will be accomplished by expanding the interval fields of the Job to be able to handle more complicated combinations of job scheduling. The current design only allows for specific interval units: Second, Minute, Hour, Day, Week, Month, and Year. Furthermore, these are strictly relative to the start date. The new design will allow more elaborate units, such as specific days of the week or month, or year, as well as combinations of these, and they would be tied to the calendar, not the start date. This will be accomplished by substituting the current design’s intervalUnit use of Calendar constants with a hierarchy of DateUnit classes, some of whom are marker extensions only: Second, Minute, Hour, Day, Week, Month, Year, DaysOfWeek, WeekOfMonth, DayOfMonth, WeekMonthOfYear, and DayOfYear. The processor will be responsible for knowing how to interpret this.

Second, Minute, Hour, Day, Week, Month, Year: Marker extension of DateUnit that act in the same manner as in the current design’s Calendar constants. They simply represent an specific amount of time to wait before the next job run.

DaysOfWeek: Defines specific days of the week when the Job is to be run. It can define one to seven days of the week. There will exist two convenience implementations that encompass weekdays and weekends.

WeekOfMonth: Defines a specific day of the week in a month. For example, one can define the first Saturday of a month. This might be ideal for setting up notifications for a meeting that must occur on a specific day of the business week but also at the start of a month.

DayOfMonth: Defines a specific day in a month. For example, one can define the 15th day of a month. This could be used for monthly check generation for disgruntled TopCoder designers.

WeekMonthOfYear: Defines a specific day of the week in a month in a given year. For example, one can define the first Saturday of January. This might be ideal for post-new year's celebrations pink slip generation for rowdy employees.

DayOfYear: Defines a specific day in a year. For example, one can define the 15th day of a year. This could be similar to the WeekMonthOfYear date unit, except when the date into the year must be consistent regardless of the vagaries of the calendar with its leap years.

Apart from date units, it is important to also incorporate the idea of intervals and recurrences. The first item is already incorporated in the current design, and it allows for the schedule to occur every number of date units. For example, we might want something to occur on the first Saturday of a month, but every second month.

The other concept of recurrence simply means that we might want a Job to run a specified amount of times. This might work in lieu of a end date, or with an end date to indicate something like: "Do this on the Saturday of every second week five times or until March 31st."

3) *CRUD functionality*

The current scheduler already incorporates CRUD operations, so this will not change fundamentally. What will change is that the CRUD operations will occur via an interface. The scheduler itself will not be affected, but the processor that will use it will not be tied to a specific implementation, as it will be accessed via a interface using Strategy, but that is more in the scope of the processor design.

4) *Current design*

It is very much desired to keep the changes to the code as few as possible. Fortunately, this is very possible here. The Job class will mostly not change, except in that the protected methods will be made public, so the processor can more easily work with them, and the interval information, as mentioned in section

3 above, will be enhanced. Most of the configuration will not change either. The link to the scheduler will be cut, though, and this will have some consequences.

In fact, only the Scheduler will be modified on a large scale. A new Scheduler interface will be created, and two implementations of it will be added, one that will read from ConfigManager as is done currently (it's writing ability will be left intact), and one that uses the Informix database. These scheduler implementations will have no interaction with the removed processor, and the methods to manage the running of jobs will be moved to the processor, such as start(), stop(), shutdown(), getJobExecutionList(), stopJob(), etc.

One of the aspects of the Job was the addition of the TriggerEventHandler. This, unfortunately, had a reference to the processor required to run it. As such, this event handler will be added by the processor at the time it reads the version 2.0 jobs. Also, Email alerts are fully configured independently of the Scheduler.

One final note about the persistence operations. When a Job or JobGroup is retrieved, it will contain the associated JobGroups and Jobs, respectively. But these will contain just the name of each, and not a recursive relationship. If the user wants to know more about them, they will have to retrieve them explicitly. This equally applies to creating, updating and deleting. When creating and updating a Job, for instance, the JobGroups must already exist, and will not be created. When deleting a Job, the JobGroups will not be deleted as well. This applies to CRU operations for JobGroups equally.

1.1 Design Patterns

The EventHandler uses the **listener** pattern. It is used to send email alerts or trigger the dependent job to start in this version.

The Scheduler implementations will be used as **Strategies** by other components.

This component uses the **Type Safe** pattern to define the Job Types, but these days, this does not really apply as anything but a fundamental pattern.

1.2 Industry Standards

None.

1.3 Required Algorithms

1.3.1 Existing algorithms

This section will elaborate existing algorithms.

Implementation:

When a job executes successfully, fails or even doesn't start, the job will raise the corresponding event. The EmailEventHandlers are listening to the job and will send email alerts according to the event.

Here is the pseudo code:

```
if (event.equals(requiredEvent)) {
//1. Generate the message
NodeList msgdata = job.getMessageData();
Node[] nodes;
if (msgdata == null) nodes = new Node[2];
else nodes = new Node[msgdata.getNodes().size() + 2];

nodes[0] = new Field("JobName", job.getName(), null, true);
nodes[1] = new Field("JobStatus", event, null, true);
copy the nodes in msgdata.getNodes() to the rest of the array.

TemplateFields data = new TemplateFields(nodes, template);
String message = DocumentGenerator.getInstance().applyTemplate(data);

// 2. send the email using EmailEngine.

    TCSEmailMessage email = new TCSEmailMessage();
    email.setSubject(emailAlertSubject);
    email.setFromAddress(emailFromAddress);
    for each recipient in recipients {
        email.addToAddress(recipient, 0);
    }
    EmailEngine.send(email);
}
```

1.3.2 *Logging*

Logging is performed in three places in this component: DBScheduler, ConfigManagerScheduler, and EmailEventHandler. This section serves as an authoritative list of where and how logging should be performed, in lieu of putting this information in method documentation.

The entry and exit of all methods will be logged at TRACE level. Any exception is logged at ERROR level.

1.4 **Component Class Overview**

Scheduler

This interface defines the CRUD operations for Jobs and JobGroups. It defines two extensions: One for reading-only data from the ConfigManager, and the second for reading from and writing to a database.

Job

This class is the specific job instance to schedule. A Job must have either a scheduled date time or a task dependency but not both.

JobType



An enumeration of the two types of Jobs: External and Java Class. This class extends the Enum class from the Type Safe Enumeration component.

JobGroup

A job group is a grouping of jobs to simplify job scheduler configuration. A group can contain one or more jobs. A job can belong to one or more groups. In current version, the job groups are used to ease "Alert Notification Configuration".

Dependence

A job can be dependent on another job regarding execution time. This class represents this relationship. It has a dependent job name, dependent event and delay properties.

EventHandler

There are three event of a job: Not Started, Executed Successful, and Executed Failed. The EventHandler is designed to handle these events of jobs.

EmailEventHandler

It is used to send email alert notifications when a job raises an event matching the requiredEvent initialized in the constructor. The email message is generated by Document Generator component.

ScheduledEnable

The ScheduledEnable interface is the only required interface for classes run as a job in Job Scheduler 2.0. It extends Schedulable interface (v1.0) and Runnable interface, and forces the implementation classes to return running status and message data at runtime.

Schedulable

The Schedulabe interface is one of two required interfaces (the other being Runnable) for classes to be run as a job.

DateUnit

A marker interface representing a unit of time to be used in configuring the schedule of the Job. Many of the implementations will refer to simply intervals, such as every second, day, week, or so, based on the start date and time. Some will be more complex, however, like referring to a specific day of the year or month. The Job will use this to indicate a specific date or time to do the job, or a



date or time from the start date to perform it. It will work in conjunction with the interval and recurrence to track when and how often the job is done.

Second

A marker extension of DateUnit that represents a base interval unit of a second. Jobs configured with it will be run in intervals of seconds, depending on the associated interval setting.

Minute

A marker extension of DateUnit that represents a base interval unit of a minute. Jobs configured with it will be run in intervals of minutes, depending on the associated interval setting.

Hour

A marker extension of DateUnit that represents a base interval unit of an hour. Jobs configured with it will be run in intervals of hours, depending on the associated interval setting.

Day

A marker extension of DateUnit that represents a base interval unit of a day. Jobs configured with it will be run in intervals of days, depending on the associated interval setting.

Week

A marker extension of DateUnit that represents a base interval unit of a week. Jobs configured with it will be run in intervals of weeks, depending on the associated interval setting.

Month

A marker extension of DateUnit that represents a base interval unit of a month. Jobs configured with it will be run in intervals of months, depending on the associated interval setting.

Year

A marker extension of DateUnit that represents a base interval unit of a year. Jobs configured with it will be run in intervals of years, depending on the associated interval setting.



DaysOfWeek

A marker extension of DateUnit that defines specific days of the week when the Job is to be run. It can define one to seven days of the week. There will exist two convenience implementations that encompass weekdays and weekends.

WeekOfMonth

A marker extension of DateUnit that defines a specific day of the week in a month. For example, one can define the first Saturday of a month. This might be ideal for setting up notifications for a meeting that must occur on a specific day of the business week but also at the start of a month.

DayOfMonth

A marker extension of DateUnit that defines a specific day in a month. For example, one can define the 15th day of a month. This could be used for monthly check generation for disgruntled TopCoder designers.

WeekMonthOfYear

A marker extension of DateUnit that defines a specific day of the week in a month in a given year. For example, one can define the first Saturday of January. This might be ideal for post-new year's celebrations pink slip generation for rowdy employees.

DayOfYear

A marker extension of DateUnit that defines a specific day in a year. For example, one can define the 15th day of a year. This could be similar to the WeekMonthOfYear date unit, except when the date into the year must be consistent regardless of the vagaries of the calendar with its leap years.

ConfigManagerScheduler

A Scheduler implementation that uses the ConfigManager as the persistent data source. As in version 2.0, it will perform all CRUD operations.

DBScheduler

A Scheduler implementation that uses the database as the persistent data source. It works with all CRUD methods. It uses connections obtained from the DBConnectionFactory for access and generate ids using the IDGenerator component.

1.5 Component Exception Definitions

This component defines one custom exception

SchedulingException

This exception will be thrown when a Scheduler operation fails. This will pertain strictly to persistence operations on Jobs or JobGroups.

ConfigurationException

This exception will be thrown when there is a configuration problem in the scheduler constructors.

1.6 Other Design Considerations

1.6.1 *Configuration File Data Items*

The Property Name Space is the job name.

Within each property name space the following name:value pairs are stored:

StartDTG: The start date / time for the job.

EndDTG: The end date/time for the job.

IntervalValue: The numeric interval value for the job.

IntervalUnit: The unit of time (see #4 above) for the interval value.

JobType: Whether the job is an external command or internal java class.

JobCmd: The job execution string (either external command or java class name, depending on the JobType value).

1.6.2 *Invalid Jobs*

The following conditions would make a job invalid during creation or modification:

- a. A start date past the end date.
- b. A name which already exists.
- c. A name which is null or an empty String.
- d. A negative or 0 increment value.
- e. An increment unit value not following one of the DataUnit types outlined in the introduction
- f. A job type not one specified in the Scheduler Object.
- g. A null start date.

1.7 Thread Safety

This component is almost thread-safe. Scheduler, JobGroup, Dependence, and EmailEventHandler class are thread-safe. The only exception is Job class; some properties in it are mutable, so the user can make modifications and persist them. However, it is not anticipated that the same Job will be handled by more than one thread.



Also, the implementations of the `Schedulable` or `ScheduledEnable` interfaces should be thread-safe, that means when the job is running, the `isDone`, `getRunningStatus`, etc methods can be called in thread-safe way.

2. Environment Requirements

2.1 Environment

- . Development language: Java1.4
Compile target: Java1.4, Java1.5

2.2 TopCoder Software Components:

Configuration Manager 2.1.5

Used for text-based persistence of a schedule, as well as the configuration of the DB scheduler.

Document Generator 2.0

Used to generate the email message.

Email Engine 3.0

Used to send the email.

TypeSafe Enum 1.0

Provides type safe enum functionality for this pre Java 5 component.

Base Exception 1.0

Provides a uniform base exception class.

DB Connection Factory 1.0

Provides connections to the `DBScheduler` class.

ID Generator 3.0

Provides ID generation for the `DBScheduler` class.

Executable Wrapper 1.0

Provides access to command line execution for an external Job.

Logging Wrapper 2.0

Used by `DBScheduler`, `ConfigManagerScheduler`, and `EmailEventHandler` to log.

Note: The Alert Factory and Event Email Processor are not used. Because both of them need to persist the messages and are a little complex. In this component, it is not worth to add this complexity while using these two components.

2.3 Third Party Components:

None.

3. Installation and Configuration

3.1 Package Name

`com.topcoder.util.scheduler.scheduling`

3.2 Configuration Parameters

3.2.1 ConfigManagerScheduler Configuration Parameters

| Parameter | Description | Values |
|---|---|--|
| <JobName> | Define a job | Can be any string except "DefinedGroups" and "Logger" |
| <JobName>.StartDate | The start date of the Job Required if this job is not dependent on another job. | A valid date |
| <JobName>.StartTime | The start time of the Job Required if this job is not dependent on another job. | A valid time of day, in milliseconds |
| <JobName>.EndDate | The end date of the job Required if this job is not dependent on another job. | A valid date, after start time |
| <JobName>.JobType | The type of the Job Required | JOB_TYPE_EXTERNAL or JOB_TYPE_JAVA_CLASS |
| <JobName>.JobCommand | The job command name Required | If external, then a command name, if java class, then the name of the class. |
| <JobName>.Active | Flag whether the job is active or not Required | True or false |
| <JobName>.Recurrence | A number stating how many time this job will be run | A positive number |
| <JobName>.Interval.Value | The value of the interval Required | Any positive number |
| <JobName>.Interval.Unit.Type | The class name of the DateUnit Required | A valid DateUnit class name |
| <JobName>.Interval.Unit.DateUnitDays | The day or days that make up the date unit. If days, these will be comma-delimited. Required if the type demands it. | A positive integer valid for the date unit type. |
| <JobName>.Interval.Unit.DateUnitWeek | The week that makes up the date unit. Required if the type demands it. | A positive integer valid for the date unit type. |
| <JobName>.Interval.Unit.DateUnitMonth | The month that makes up the date unit. Required if the type demands it. | A positive integer valid for the date unit type. |
| <JobName>.ModificationDate | The date this Job was last modified. Required | A date |
| .<JobName>.Dependence | Specify the depended job, Optional | |
| <JobName>.Dependence.<dependedJobName> | The name of the depended job | Must have been defined in the configuration file |
| <JobName>.Dependence.<dependedJobName>.Status | Tell when the job should be started, on which status of the completion of the depended job, required | One of SUCCESSFUL, FAILED or BOTH |
| <JobName>.Dependence.<dependedJobName>.Delay | Specify a time delay before the job is triggered, optional. | Non-negative value, the unit is ms. |
| <JobName>.Messages | Define the message alerts of this job, Optional | |

| | | |
|--|--|---|
| <JobName>.Messages.<Status> | When to send the email alert, optional | One of SUCCESSFUL, FAILED and NOTSTARTED |
| <JobName>.Messages.<Status>.TemplateText | The name of the template text to generate message, optional. If it is not specified, the default template in scheduler will be used. | File name |
| <JobName>.Messages.<Status>.FromAddress | The From Email Address of the message alerts, required | Email address |
| <JobName>.Messages.<Status>.Subject | The Subject of the message alerts, required | Subject |
| <JobName>.Messages.<Status>.Recipients | The Recipients of the message alerts, required | Email addresses |
| DefinedGroups | Define job groups, optional | |
| DefinedGroups.<GroupName> | The job group name, optional | |
| DefinedGroups.<GroupName>.Jobs | The jobs this group contained, required, at least one. | The name of defined jobs |
| DefinedGroups.<GroupName>.Messages | The message alerts configuration based on the group, similar with configuration on a single job. Optional | Similar with the configuration on a single job. |
| Logger | Name of logger | Optional. Defaults to fully-qualified name of ConfigManagerScheduler class. |

3.2.2 DBScheduler Configuration Parameters

| Parameter | Description | Values |
|----------------------------|---|---|
| ConnectionFactoryClassName | Fully-qualified name of the DBConnection Factory | "com.topcoder.db.connectionfactory.DBConnectionFactoryImpl" Optional. Will default to the value above. |
| ConnectionFactoryNamespace | Namespace that the DBConnection Factory implementation uses. | "test" Required. |
| ConnectionName | Name of the connection to the persistence to get from the DBConnection Factory | "PersistenceConnection" Required. |
| IDGenSeqName | Named sequence for the tables. Used to retrieve an IDGenerator that can service it. | "SchedulerId" Required. |
| IDGenClassName | Name of the IDGenerator class that services the named sequence | "com.topcoder.util.idgenerator.OracleSequenceGenerator" Optional. Will attempt to find a generator already configured to handle the name sequence. |
| Logger | Name of logger | Optional. Defaults to fully-qualified name of DBScheduler class. |

3.3 Dependencies Configuration

The developer should refer to the component specification of the used TopCoder components to configure them.

3.3.1 DDL for DBScheduler

```
CREATE TABLE JOB (  
  JobId INT NOT NULL PRIMARY KEY,  
  Name VARCHAR(40) NOT NULL,  
  StartDate DATE,  
  StartTime INT,  
  EndDate DATE,  
  DateUnit VARCHAR(60) NOT NULL,  
  DateUnitDays VARCHAR(20), // this one will be comma-delimited  
  DateUnitWeek VARCHAR(2),  
  DateUnitMonth VARCHAR(2),  
  Interval INTEGER NOT NULL,  
  Recurrence INTEGER NOT NULL,  
  Active CHAR(1) NOT NULL,  
  JobType VARCHAR(20) NOT NULL,  
  JobCommand VARCHAR(40) NOT NULL,  
  DependenceJobName VARCHAR(60),  
  DependenceJobStatus VARCHAR(20),  
  DependenceJobDelay VARCHAR(20)  
);
```

```
CREATE TABLE Message (  
  MessageId INT NOT NULL PRIMARY KEY,  
  OwnerId INT NOT NULL,  
  Name VARCHAR(40) NOT NULL,  
  FromAddress VARCHAR(40) NOT NULL,  
  Subject VARCHAR(40) NOT NULL,  
  TemplateText VARCHAR(40),  
  Recipients VARCHAR(255) NOT NULL // this one will be comma-delimited  
);
```

```
CREATE TABLE Group (  
  GroupId INTEGER NOT NULL PRIMARY KEY,  
  Name VARCHAR(40)  
);
```

```
CREATE TABLE GroupJob (  
  GroupId INTEGER NOT NULL,  
  JobId INTEGER NOT NULL  
);
```

4. Usage Notes

4.1 Required steps to test the component

- Extract the component distribution.
- Follow [Dependencies Configuration](#).

- Execute 'ant test' within the directory that the distribution was extracted to.

4.2 Required steps to use the component

Follow demo.

4.3 Demo

This demo will show the use of the DBScheduler, since it works on all operations. The ConfigScheduler is worked on in the same manner, but it only has retrieve ops.

```
// Instantiate the DBScheduler, passing it the name of the config file
// containing job data.
Scheduler myScheduler = new DBScheduler(NAMESPACE);

// add new jobs.
// This job will start at 1 am on the 10th of March (GregorianCalendar months
// run from 0 to 11), and will run once a day, at 1 am, everyday until
// the 10th of March 2004 (inclusive).
Job deleteFiles = new Job("Nightly file cleanup", JobType.JOB_TYPE_EXTERNAL,
    "erase *.tmp");
deleteFiles.setStartDate(new GregorianCalendar(2003, 04, 10, 01, 00, 00));
deleteFiles.setStopDate(new GregorianCalendar(2004, 04, 10, 01, 00, 00));
deleteFiles.setIntervalUnit(new Day());
myScheduler.addJob(deleteFiles);

// Add a job dependent on another job regarding the execution time.
// jobA has a specific schedule time. jobB is dependent on jobA and
// jobC is dependent on jobB
Job jobA = new Job("jobA", JobType.JOB_TYPE_JAVA_CLASS,
    "com.topcoder.util.scheduler.scheduling.MyJob");
// Calendar representing the date
jobA.setStartDate(new GregorianCalendar(2003, 04, 10, 01, 00, 00));
// long representing a time of day
jobA.setStartTime(580);
// Calendar representing the date
jobA.setStopDate(new GregorianCalendar(2004, 04, 10, 01, 00, 00));
jobA.setIntervalUnit(new Day());
jobA.setIntervalValue(5);

myScheduler.addJob(jobA);

// another job with name jobB
Job jobB = new Job("jobB", JobType.JOB_TYPE_EXTERNAL, "dir");
// the delay is 10s
jobB.setDependence(new Dependence("jobA", EventHandler.SUCCESSFUL, 10000));
jobB.setIntervalUnit(new Week());
jobB.setIntervalValue(1);

myScheduler.addJob(jobB);

// another job, dependent on jobB, and configured for no delay
Job jobC = new Job("jobC", JobType.JOB_TYPE_EXTERNAL, "java -version");
jobB.setDependence(new Dependence("jobB", EventHandler.SUCCESSFUL, 0));
jobC.setIntervalUnit(new Month());
jobC.setIntervalValue(5);

// add email alert event handler to jobC, if the jobC executed
// unsuccessfully, an email will be sent to name1@topcoder.com
// the typical messageTemplate is like
//
// The Job %JobName% is %JobStatus%...
//
Log logger = LoggerFactory.getLog();
Template template = new XsltTemplate();
template.setTemplate("The Job %JobName% is %JobStatus%...");
EmailEventHandler handler1 = new EmailEventHandler(template, Arrays.asList(new
String[] { "name1@topcoder.com",
    "service@topcoder.com", "Failure of Job" } ), EventHandler.FAILED,
```



```
"admin@topcoder.com", "Notification", logger);
jobC.addHandler(handler1);

myScheduler.addJob(jobC);

// add a job group to scheduler
JobGroup group = new JobGroup("group_1", Arrays.asList(new Job[] {jobA, jobB,
jobC}));

myScheduler.addGroup(group);

// add an email Event Handler to the group
// the following code means if any one of the jobs in the group executed
// successfully, an email alert will be sent to "name2@topcoder.com" and
// "name3@topcoder.com"
EmailEventHandler handler2 = new EmailEventHandler(template, Arrays.asList(new
String[] {"name2@topcoder.com",
"name3@topcoder.com", "service@topcoder.com", "Success of Job"}),
EventHandler.SUCCESSFUL, "admin@topcoder.com", "Notification", logger);
group.addHandler(handler2);

// you can remove the handler from job and group, but you then have to
// update each
jobC.removeHandler(handler1);
group.removeHandler(handler2);

myScheduler.updateJob(jobC);
myScheduler.updateGroup(group);

// you also can remove job and group from the scheduler
myScheduler.deleteGroup(group.getName());
myScheduler.deleteJob(jobA);
```

5. Future Enhancements

Add more EventHandlers to provide more functions in Job Scheduler.

.